

Education

Ph.D. Student in Computer Science

Rice University

Advisor: Dr. Yuke Wang

2025–May 2030 (Expected)

Houston, TX, USA

Ph.D. Student in Computer Engineering

Indiana University

Advisor: Dr. Dingwen Tao

2023–2025

Bloomington, IN, USA

B.S. in Physics

University of Science and Technology of China

Advisor: Dr. Changling Zou

2019–2023

Hefei, Anhui, China

Research Interests

- Efficient systems for generative AI
- System optimization for multimodal LLMs
- Data compression and communication in HPC/ML training

Publications

- **Preprint** Fanjiang Ye, Zepeng Zhao, Yi Mu, Jucheng Shen, Renjie Li, Kaijian Wang, Desen Sun, Saurabh Agarwal, Myungjin Lee, Triston Cao, Aditya Akella, Arvind Krishnamurthy, T. S. Eugene Ng, Zhengzhong Tu, Yuke Wang. **SUPERGEN: An Efficient Ultra-high-resolution Video Generation System with Sketching and Tiling.** [arXiv](#)
Video Generation : Built SUPERGEN, a training-free tile-based framework with region-aware caching and communication minimized multi-GPU tile parallelism for efficient, high-quality ultra-high-resolution video generation.
- **NeurIPS'25** Jinda Jia, Cong Xie, Hanlin Lu, **Fanjiang Ye**, Hao Feng, Daoce Wang, Haibin Lin, Zhi Zhang, Xin Liu. **DUO: No Compromise to Accuracy Degradation.** [To Appear](#)
Communication Compression in Distributed LLM : Introduced DUO: overlaps an extra high-precision gradient sync within compute to hide communication and recover accuracy under aggressive gradient compression.
- **ICML'25 Spotlight** Xiyuan Wei, Ming Lin, **Fanjiang Ye**, Fengguang Song, Liangliang Cao, My T. Thai, Tianbao Yang. **Model Steering: Learning with a Reference Model Improves Generalization Bounds and Scaling Laws.** [arXiv](#)
CLIP Optimization : Formalized model steering (DRRho/DRO) and introduced DRRho-CLIP for reference-guided training with better generalization, data efficiency, and scaling.
- **ICS'25 Best Paper Runner-up** Boyuan Zhang, Bo Fang, **Fanjiang Ye**, Luanzheng Guo, Fengguang Song, Tallent Nathan, Dingwen Tao. **BMQSim: Overcoming Memory Constraints in Quantum Circuit Simulation with a High-Fidelity Compression Framework.** [arXiv](#)
Compression in Quantum Computing : Designed BMQSim, a compression-aware quantum circuit simulator with GPU-based lossy compression, circuit partitioning, pipeline-integrated data movement, and two-level memory management.
- **SC'24** Hao Feng, Boyuan Zhang, **Fanjiang Ye**, Min Si, Ching-Hsiang Chu, Jiannan Tian, Chunxing Yin, Zhaoxia (Summer) Deng, Yuchen Hao, Pavan Balaji, Tong Geng, Dingwen Tao. **Accelerating Communication in Deep Learning Recommendation Model Training with Dual-Level Adaptive Lossy Compression.** [doi](#)
Compression in DLRM : Accelerated DLRM with error-bounded compression for embedding all-to-all, via dual-level adaptive bounds and GPU-optimized tensors.
- **PPoPP'25 Poster** Boyuan Zhang, Luanzheng Guo, Jiannan Tian, Jinyang Liu, Daoce Wang, **Fanjiang Ye**, Chengming Zhang, Jan Strube, Nathan R. Tallent, Dingwen Tao. **High-performance Visual Semantics Compression for AI-Driven Science.** [doi](#)
Compression in AI Science : Developed ViSemZ, a high-performance AI-based scientific-image compressor that preserves visual semantics via sparse encoding with variable-length truncation and optimized lossless coding.
- **Preprint** Xinrui Zhong, Xinze Feng, Jingwei Zuo, **Fanjiang Ye**, Yi Mu, Junfeng Guo, Heng Huang, Myungjin Lee, Yuke Wang. **An Efficient and Adaptive Watermark Detection System with Tile-based Error Correction.** [arXiv](#)
Diffusion Watermarking : Designed QRMark, an adaptive tile-based watermark detector with QR code error correction

and resource-aware GPU scheduling for efficient, robust large-scale detection.

- [Preprint](#) Xiyuan Wei, [Fanjiang Ye](#), Ori Yonay, Xingyu Chen, Dingwen Tao, Tianbao Yang. [FastCLIP: A Suite of Optimization Techniques to Accelerate CLIP Training with Limited Resources](#). [arXiv](#)

CLIP Optimization : Engineered FastCLIP, a distributed CLIP training framework that leverages compositional optimization and comm.-efficient gradient reduction for efficient training on limited resources.

Research Experience

Rice University, [Yuke's Laboratory](#)

Graduate Research Assistant

2025–Present

Houston, TX, USA

- Research in designing efficient system techniques for image/video generation.
- Exploring in heterogeneous and high-performance MLLM serving system.

Indiana University, [HiPDAC Laboratory](#)

Graduate Research Assistant

2023–2025

Bloomington, IN, USA

- Research in designing accelerator-based lossy compression for HPC/ML applications.
- Developing the efficient distributed CLIP training framework.

Coding Language

- Python
- C/C++
- CUDA

Professional Service

- ASPLOS'26 Artifact Evaluation Committee
- SOSP'25 Artifact Evaluation Committee
- QCE'24 Sub-Reviewer

Honors and Awards

- Indiana University Travel Awards (\$1500), Indiana University Bloomington 2024
- USTC Fellowship (\$2000), University of Science and Technology of China 2023
- Outstanding Student Scholarship (Top 25%), University of Science and Technology of China 2020, 2021, 2022

Teaching Experience

- Teaching Assistant of ENGR-E 516: Cloud Computing, Spring 2025, Indiana University